

An Introduction to IPsec

Daniel Migault

IPsec Overview

IPsec designates an architecture [[RFC4301](#)] as well as a suite of protocols AH [[RFC4302](#)], ESP [[RFC4303](#)], IKEv2 [[RFC7296](#)]

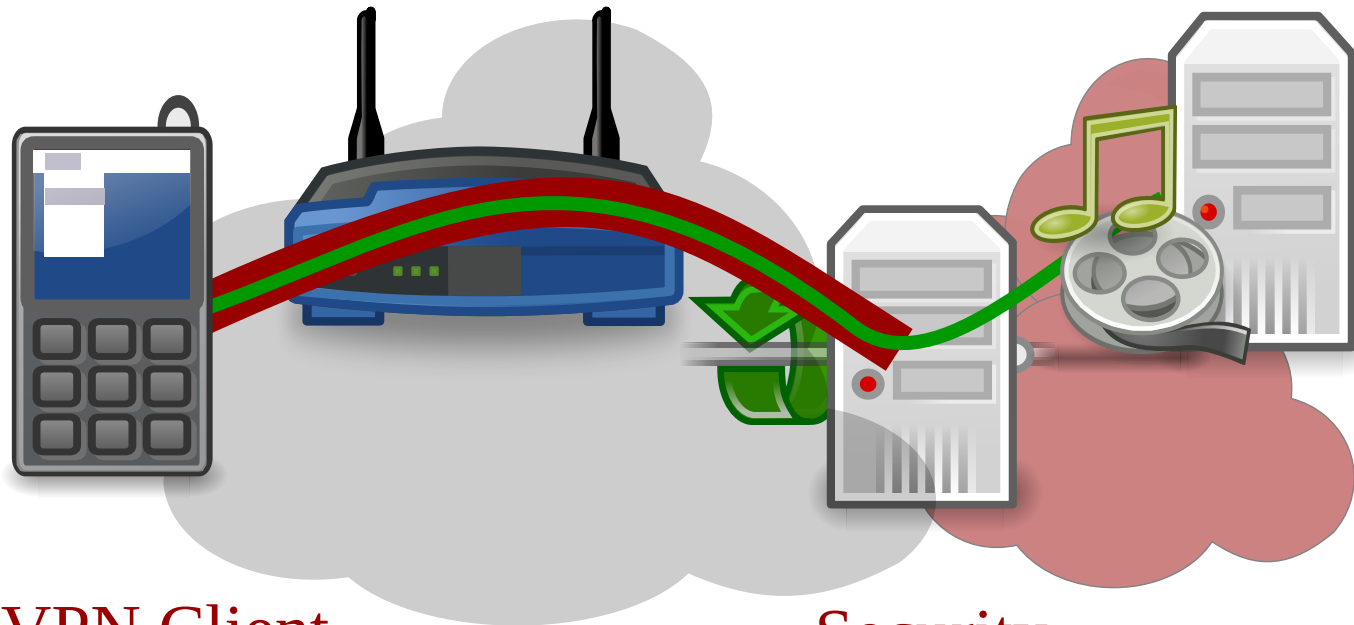
IPsec secures IP traffic for:

- Virtual Private Network (VPN):
 - Extend a trusted domain over an untrusted domain
- Gateway-to-Gateway:
 - Interconnect two trusted domain over an untrusted network
- End-to-End secure communication
 - Terminating nodes secure communications using IPsec similarly to TLS

IPsec Overview - Use Case - VPN

Untrusted
Network

Trusted
Network



VPN Client

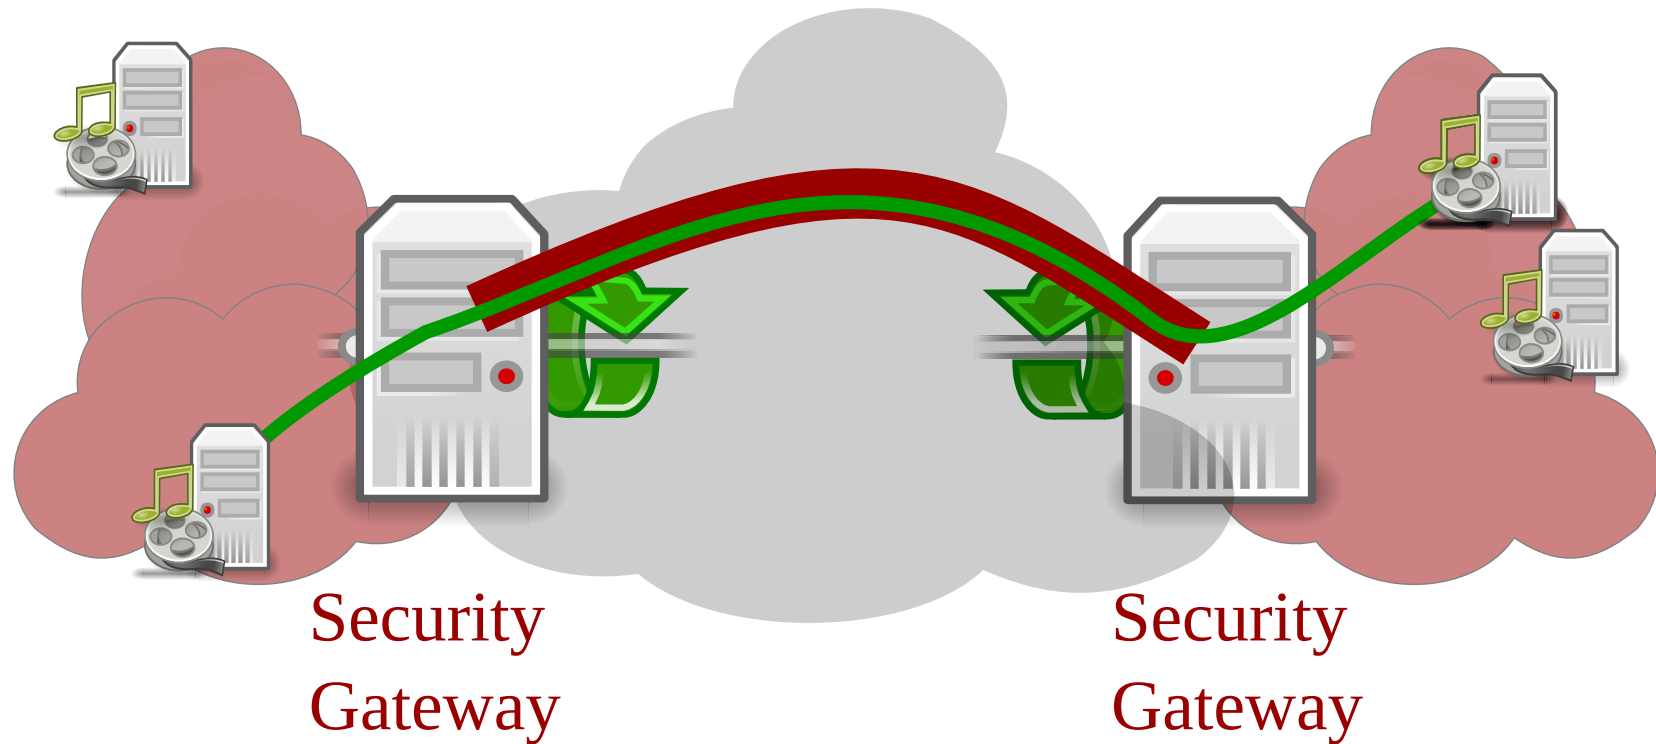
Security
Gateway

IPsec Overview - Use Case - GW-to-GW

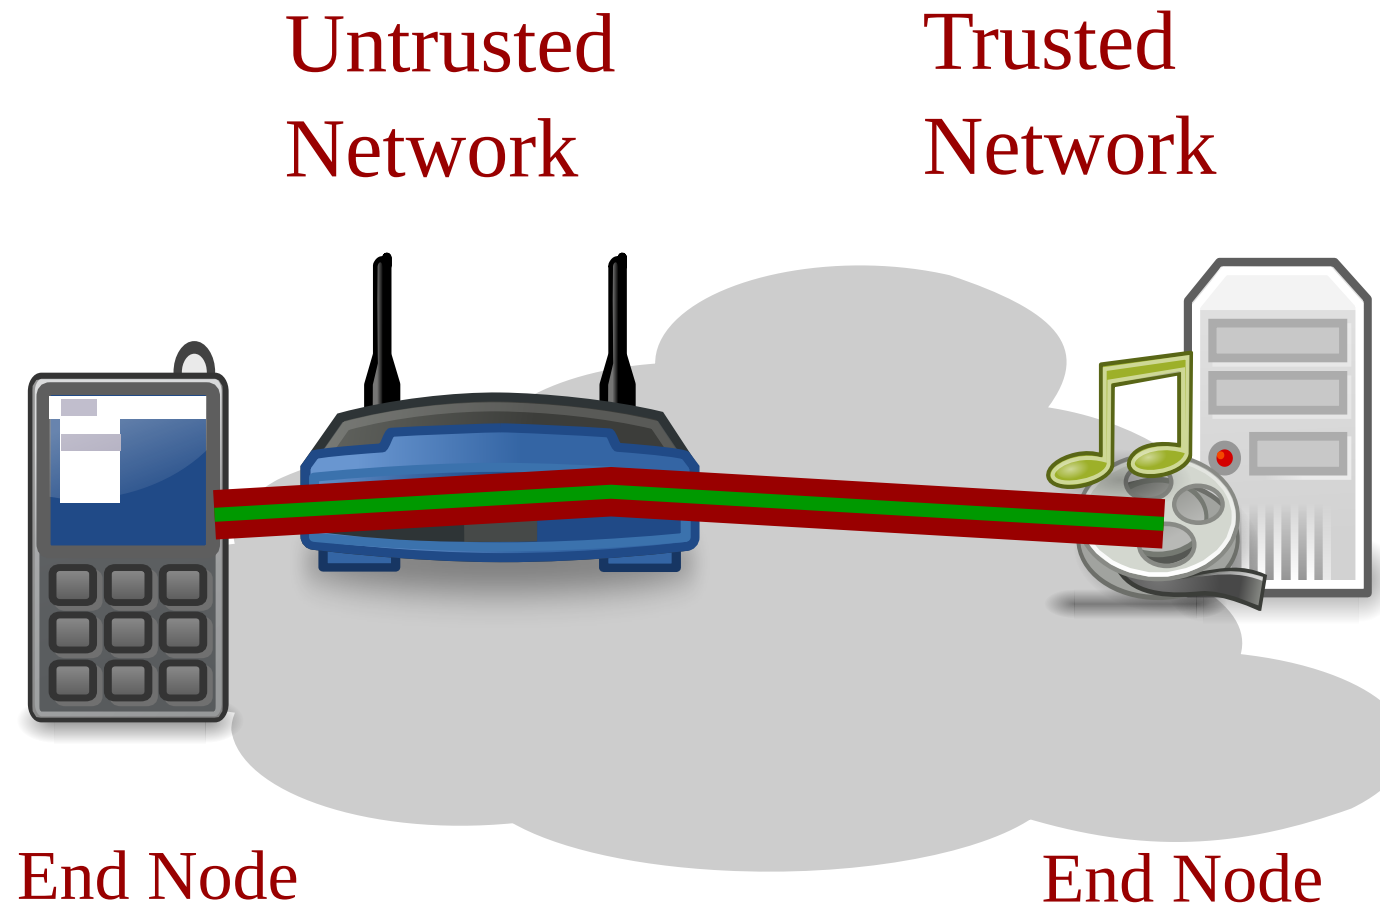
Trusted
Network

Untrusted
Network

Trusted
Network

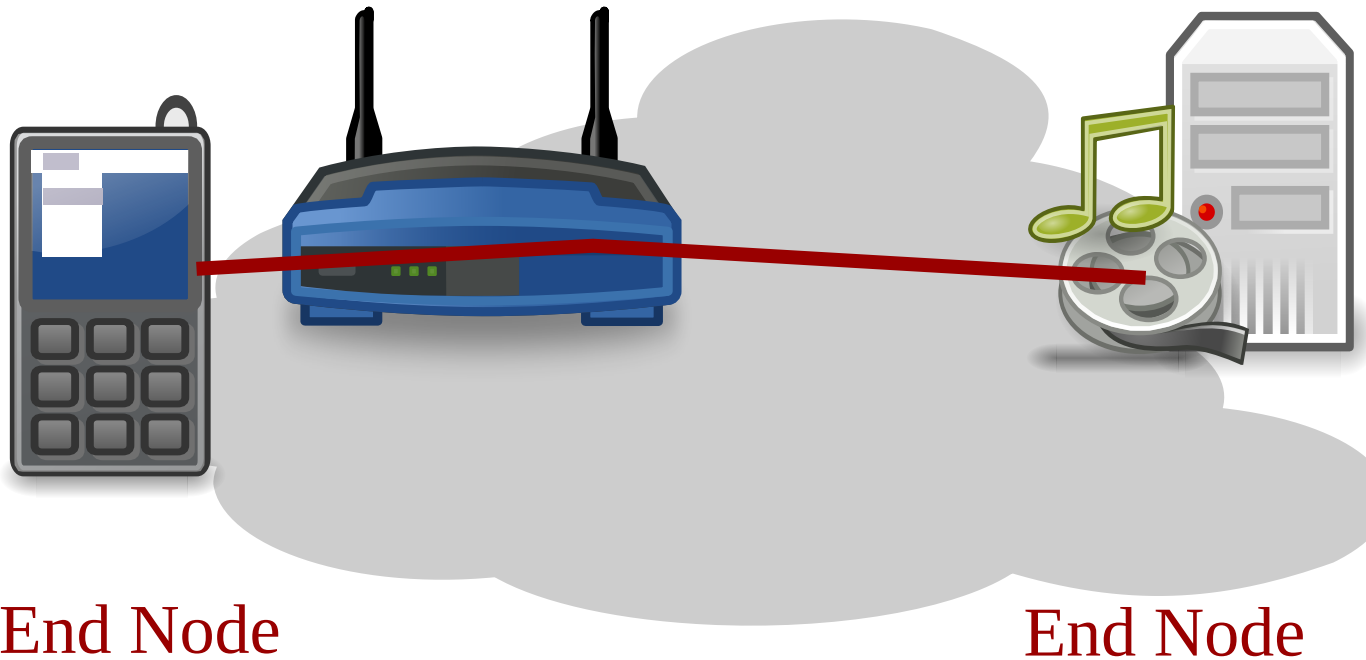


IPsec Overview - Use Case - E2E Security



IPsec Overview - Use Case - E2E Security

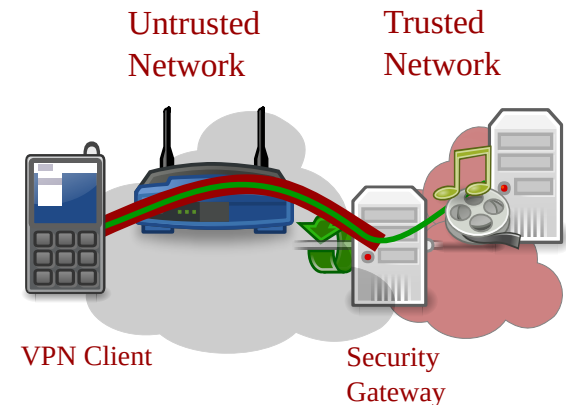
Untrusted
Network



IPsec Overview

IPsec secures traffic over untrusted network.

- Endpoints are trusted and authenticated
- Endpoint are configured to secure traffic with its peer.
 - Inbound and Outbound traffic



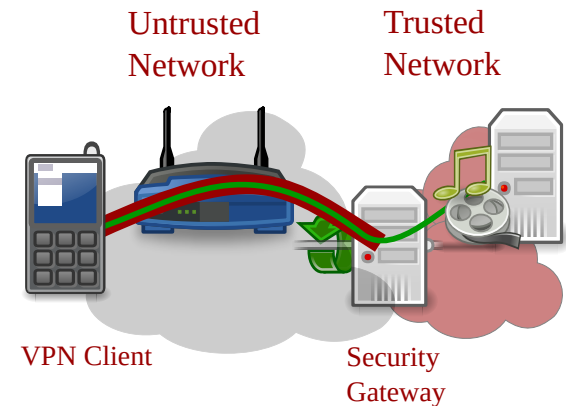
IPsec Overview

Securely exchanging traffic requires an agreement on:

- Keying material and cryptographic algorithms
- Packet format (protocols, modes, ...)

Manual configuration is possible, but using key exchange agreement is recommended.

- Internet Key Exchange Protocol (IKEv2)



IPsec Overview

Using IKEv2 enhances:

- Cryptographic material management:
 - Negotiation is delegated to the peers
 - Automated life cycle
 - ...
- Security:
 - Avoid misconfigurations (reuse of keys, IV, ...)
 - Avoids centralized key distribution and key sharing outside of the peers
 - ...

IPsec Overview - Trust Model

IPsec trust model can be recaped as follows:

- Peers trusted each other within a security domain
- Peers are configured to secure there respective inbound /outbound traffic
- Peers agree on the keying material to be used
- Peers exchange securly traffic over an untrusted network.
 - encapsulated (VPN): tunnel mode
 - protected via IP options: transport mode

IPsec Overview - Protocol Suite

IPsec defines an architecture and a suite of protocols:

- Architecture that defines the IPsec engine:
 - Security Policies (SP),
 - Security Associations (SA)
- IPsec packet processing (Data Plan):
 - ESP: Encapsulation Security *Payload*
 - AH: Authentication *Header*
- IKEv2: Internet Key Exchange Protocol (Control Plan)

IPsec Overview - Position to TLS

IPsec secures infrastructure:

- Both peers belong to the same administrative domain
 - Configuration on both peers
 - Peers authenticate each other
- IPsec secures IP traffic (IP packet, IP payload)

TLS secures Web traffic:

- TLS protects the application payloads (TCP payload)
- TLS traffic as TCP port 443 (Implicit single Security Policy)
- Only the TLS Web browser authenticates the Web Server
 - Web application authenticates the user (login, passwords)

IPsec Overview - Position to TLS

IPsec and TLS share a similar construction:

- Key Exchange protocol to configure / manage the data plan
- Exchange traffic through the data plan

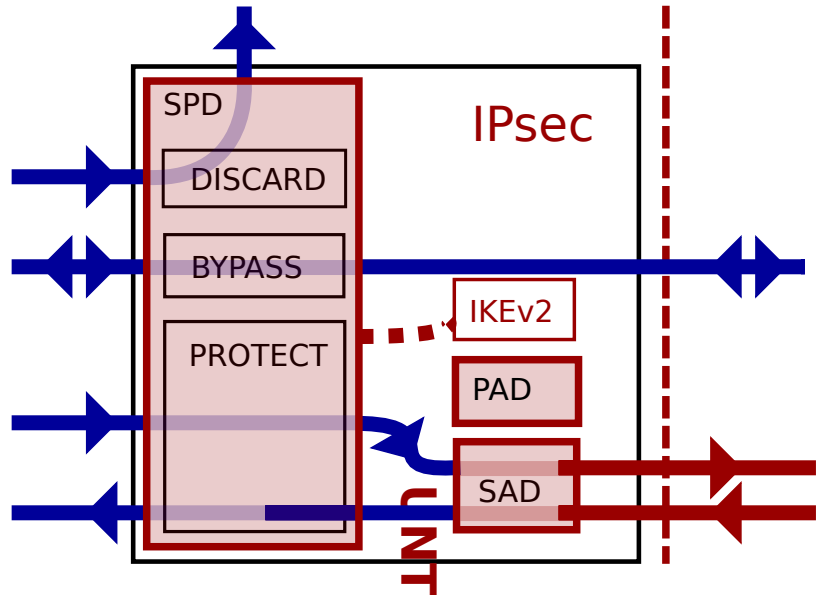
While they have different usages both protocols can likely be interchanged:

- TLS can perform mutual authentication
- IPsec can also leave one peer unauthenticated

IPsec Architecture

IPsec architecture involves:

- SPD: Security Policy Databases
- SAD: Security Association Database
- PAD: Peer Authentication Database



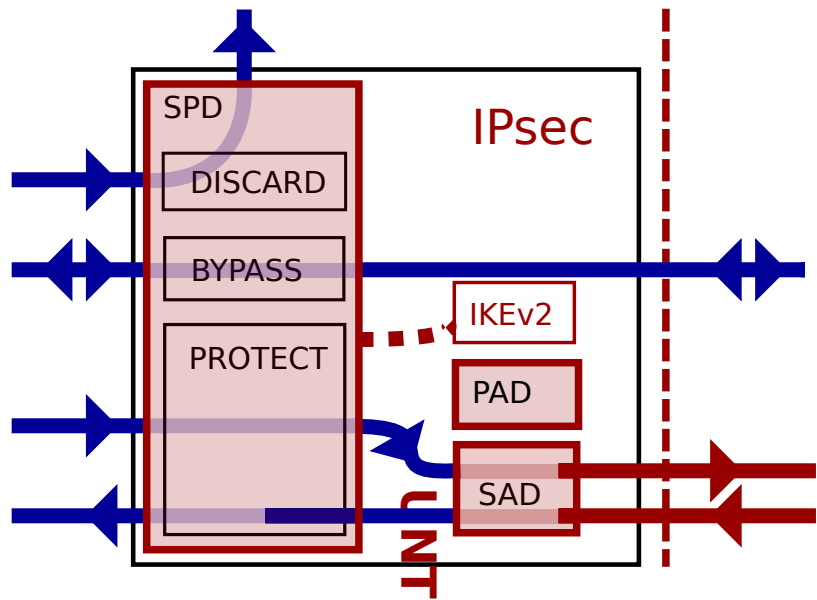
SPD provides packet policy DISCARD, BYPASS or PROTECT.
SAD provides information to encrypt and decrypt IPsec packets.
PAD provides information for IKEv2 authentication.

IPsec Architecture:

IPsec processes traffic according to these databases.

IPsec processes Inbound / Outbound differently

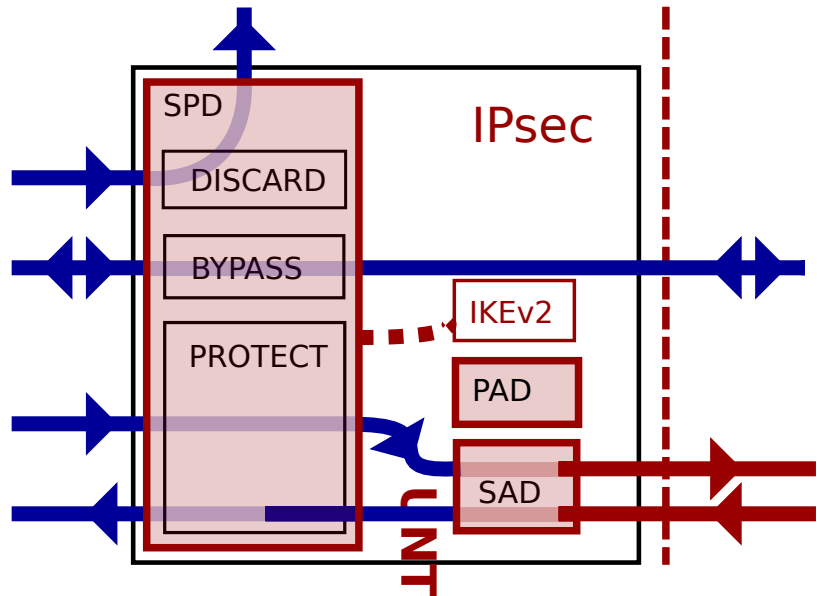
- Interactions SAD, SPD, PAD are different



IPsec Architecture - Processing Outbound Packets

Outgoing packet is *matched*
against the SPD:

- Traffic Selectors (TS)
- BYPASS, DISCARD are processed immediately
- PROTECT requires a SA



If no SA has been created:

- IKEv2 lookup the PAD to authenticate the peer
- IKEv2 proceeds to SA agreement

The packet is encrypted as indicated by the SA

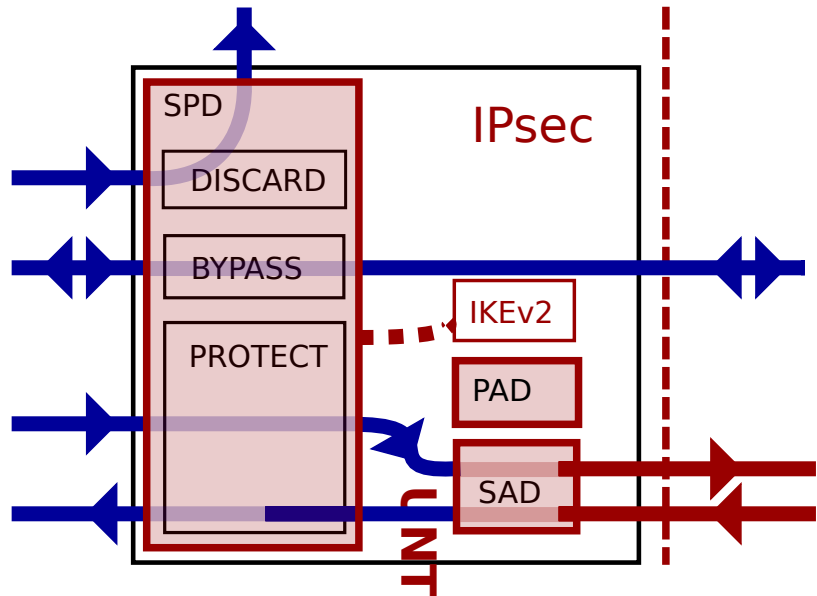
IPsec Architecture - Processing Inbound Packets

Incoming IPsec packet is *matched* against the SAD

- Security Policy Index (SPI)
- Non matching packets are discarded

SA provides information to decrypt the packet

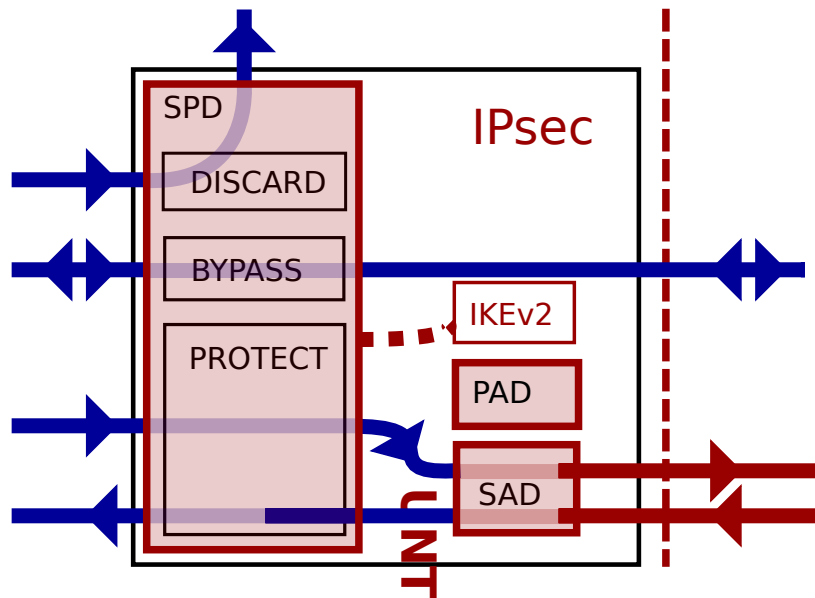
Clear text packet is matched against the SPD



IPsec Architecture - Processing Inbound Packets

Incoming non IPsec packet is *matched* against the SPD

- Traffic Selectors (TS)



IPsec Architecture - Security Policy Database

SPD contains a high level description of the Security Policies

- Matching rules are defined via Traffic Selectors (TS) that can be read from the packet (i.e IP addresses, ports...)
- Matching packet are associated a DISCARD, BYPASS or PROTECT policy

IPsec Architecture - Security Policy Database

The SPD's complexity results from opposite goals:

- Performance: matching every inbound / outbound packets
- Manageability: expressing human readable / abstract policies.

SPD is usually subdivided in multiple sub databases distributed over the system.

- SPD-I for BYPASSed inbound traffic
- SPD-O for BYPASSed outbound traffic
- SPD-S for traffic that needs to be PROTECTed
- SPD cache reflects in the kernel a packet level expression
 - derived from more abstracted expression of SP.

IPsec Architecture - Security Policy Database

SP indicates the policy associated to the traffic:

- Traffic Selectors (TS):
 - Source / Destination IP Address
 - Next Layer Protocol,
 - Source / Destination Port ...
 - Name
 - PFP Flag: Populated From Packet
- Processing Info:
 - Policy: BYPASS, DISCARD, PROTECT
 - IPsec mode, Tunnel IP header
 - IPsec protocol: ESP, AH
 - Cryptographic Algo

IPsec Architecture - Security Policy Database

SP remains abstract as TS do not necessarily match packet:

- TS uses ANY, OPAQUE words, Names, value ranges

Name designates:

- The peer ID (company.com) setting a SA with me via IKEv2
 - VPN Security Gateway (IKEv2 Responder) not the VPN Client
- Generated dynamically with system informations (UID)
 - Used by the Sender not the Destination

Special words:

- ANY: any value
- OPAQUE: field should not be looked at

IPsec Architecture - Security Policy Database

Generating SPD Cache:

- PFP defines whether:
 - Any TS match is associated to a different SA
 - All TS matches are associated to a single SA

SPD is an ordered database

SPD Cache lookup is generally faster, as such for outbound traffic
SPD Cache lookup is usually performed first, followed by a SPD lookup.

IPsec Architecture: SPD - Security Gateway Example

A Security Gateway:

- PROTECTes traffic of all the [users@company.com](#)
 - Each user is provided an IP address
 - Each user has a specific IPsec session (a different key)
- DISCARDS any other traffic

The ordered SPD will look like:

- FQDN [company.com](#) PROTECT
 - for each IP generates a specific SA
- ANY DISCARD

VPN Example: Initial SPD - SG

- BYPASSing IKEv2 channel

Traffic Selectors	PFP	SPD Cache	Processing
IP_dst: IP_{MN}^{outer}	0	IP_{MN}^{outer}	Policy: BYPASS
IP_src: IP_{SG}	0	IP_{SG}	
Port_dst: 500	0	500	
Port_src: 500	0	500	
IP_dst: IP_{SG}	0	IP_{SG}	Policy: BYPASS
IP_src: IP_{MN}^{outer}	0	IP_{MN}^{outer}	
Port_dst: 500	0	500	
Port_src: 500	0	500	

VPN Example: Initial SPD - SG

- Protecting the VPN Client traffic

Traffic Selectors	PFP	SPD Cache	Processing
Name: company.com	0	-	Policy: PROTECT
IP_dst: ANY	0	ANY	Mode: Tunnel
IP_src: ANY	1	IP_{MN}^{inner}	Tunnel Header: IP_{MN}^{outer} , IP_{SG}
Port_dst: ANY	0	ANY	Protocol: ESP
Port_src: ANY	0	ANY	Algorithm: AES-GCM

VPN Example: Initial SPD - SG

- Protecting the VPN Client traffic

Traffic Selectors	PFP	SPD Cache	Processing
Name: company.com	0	-	Policy: PROTECT
IP_dst: ANY	1	IP_{MN}^{inner}	Tunnel Header: IP_{SG} , IP_{MN}^{outer}
IP_src: ANY	0	ANY	Mode: Tunnel
Port_dst: ANY	0	ANY	Protocol: ESP
Port_src: ANY	0	ANY	Algorithm: AES-GCM

VPN Example: Initial SPD - SG

- Preventing any other traffic

Traffic Selectors	PFP	SPD Cache	Processing
IP_dst: ANY	0	ANY	Policy: DISCARD
IP_src: ANY	0	ANY	
Port_dst: ANY	0	ANY	
Port_src: ANY	0	ANY	

IPsec Architecture: SAD

SA implements the policy associated to the traffic:

- Processing info to implement the PROTECT Policy:
 - keying material and algorithm
 - IP processing information (ESP, AH, mode)
- Traffic Selectors (TS):
 - Source / Destination IP Address
 - Next Layer Protocol,
 - Source / Destination Port ...
- Security Protocol Index (SPI)

Note:

- SA TS are used to perform the SPD lookup
- SPI identifies the SA associated to encrypted packet

IPsec Architecture: SAD

Outbound packets use SA to implement a PROTECT SP:

- The appropriated SA is provided *by* the SP matching the clear text packet, using internal structure of the system (pointer)
- The SA enables encryption of the packet

Inbound packet uses the SA to implement a PROTECT SP:

- The appropriated SA is provided *by* the SPI
- The SA enables encryption of the packet

IPsec Architecture: SAD - SA Structure

The SA contains the following main information:

- Cryptographic material
 - keys and cryptographic algorithms
 - IPsec mode: Tunnel or Transport
 - Security Protocols: AH / ESP
 - SA lifetime
- IPsec signaling:
 - SPI
 - Sequence Number
 - Sequence Counter overflow
 - Anti Reaply Window
- Traffic Selector
 - inner packets

IPsec Architecture: SAD - SA Structure

The SA contains the following additional information:

- Additional checking information:
 - Stateful fragment checking
 - Don't Fragment (DF) bit
 - Differentiated Services Code Point (DSCP)
 - MTU

IPsec Architecture: PAD

Peer Authorization Database:

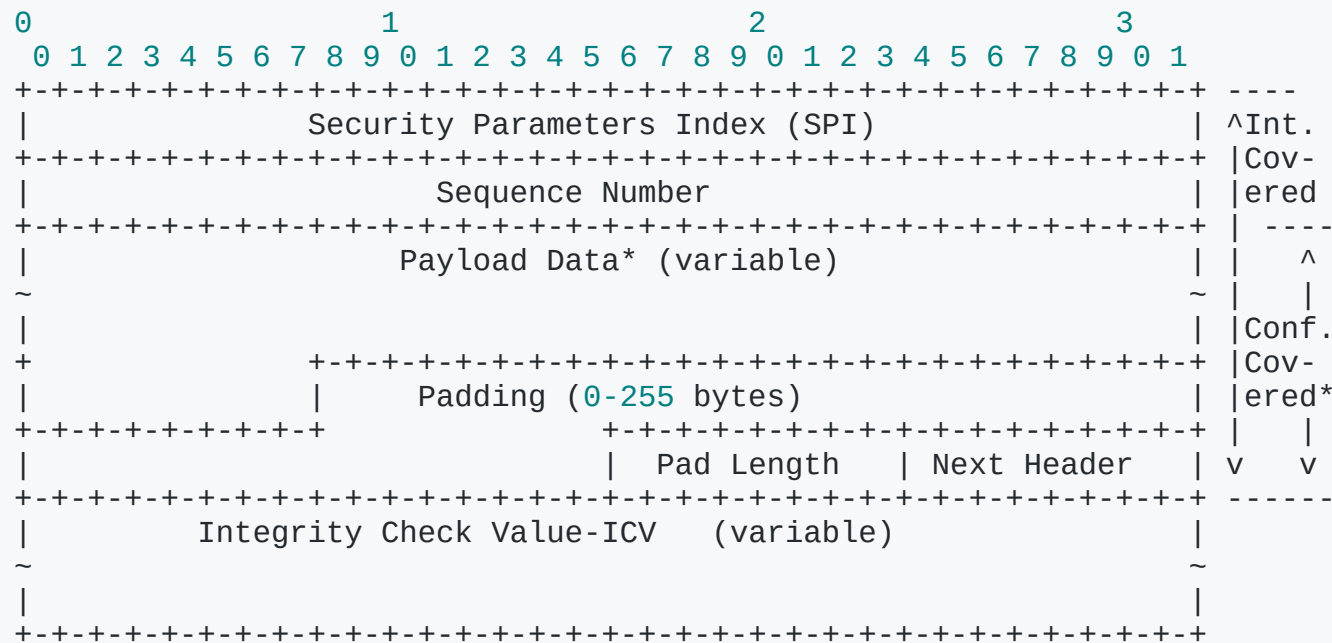
- Contains information on who can communicate with the IPsec layer and how IKEv2 must proceed to the authentication.
- Should be seen as a meta database that provides inputs for IKEv2 to configure the Security Policies.

IPsec ESP / AH

IPsec protects traffic with a combination of two protocols:

- ESP Encapsulated Security Payload securing the IP payload
 - ESP leaves unprotected the IP header
 - Widely deployed
- AH Authentication Header authenticating the IP packet
 - NAT change the IP headers
 - Hardly deployed

IPsec Encapsulated Security Payload



IPsec Encapsulated Security Payload

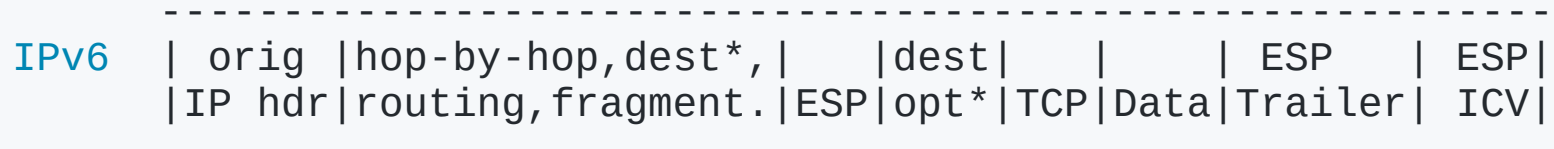
- SPI: identifies the SA for inbound packets
- Sequence Number: to provide anti-replay protection
- Payload Data: The clear text information carried by the IPsec packet
- Padding: provides 32 bit alignment
- ICV: authenticate the packet

IPsec ESP Transport mode

BEFORE APPLYING ESP



AFTER APPLYING ESP



|<--- encryption ---->|
|<----- integrity ----->|

* = **if** present, could be before ESP, after ESP, **or** both

IPsec ESP Tunnel mode

BEFORE APPLYING ESP

IPv6

```
-----  
|          | ext hdrs |          | |
| orig IP  | if present| TCP    | Data    |  
|          |          |          |          |  
-----
```

AFTER APPLYING ESP

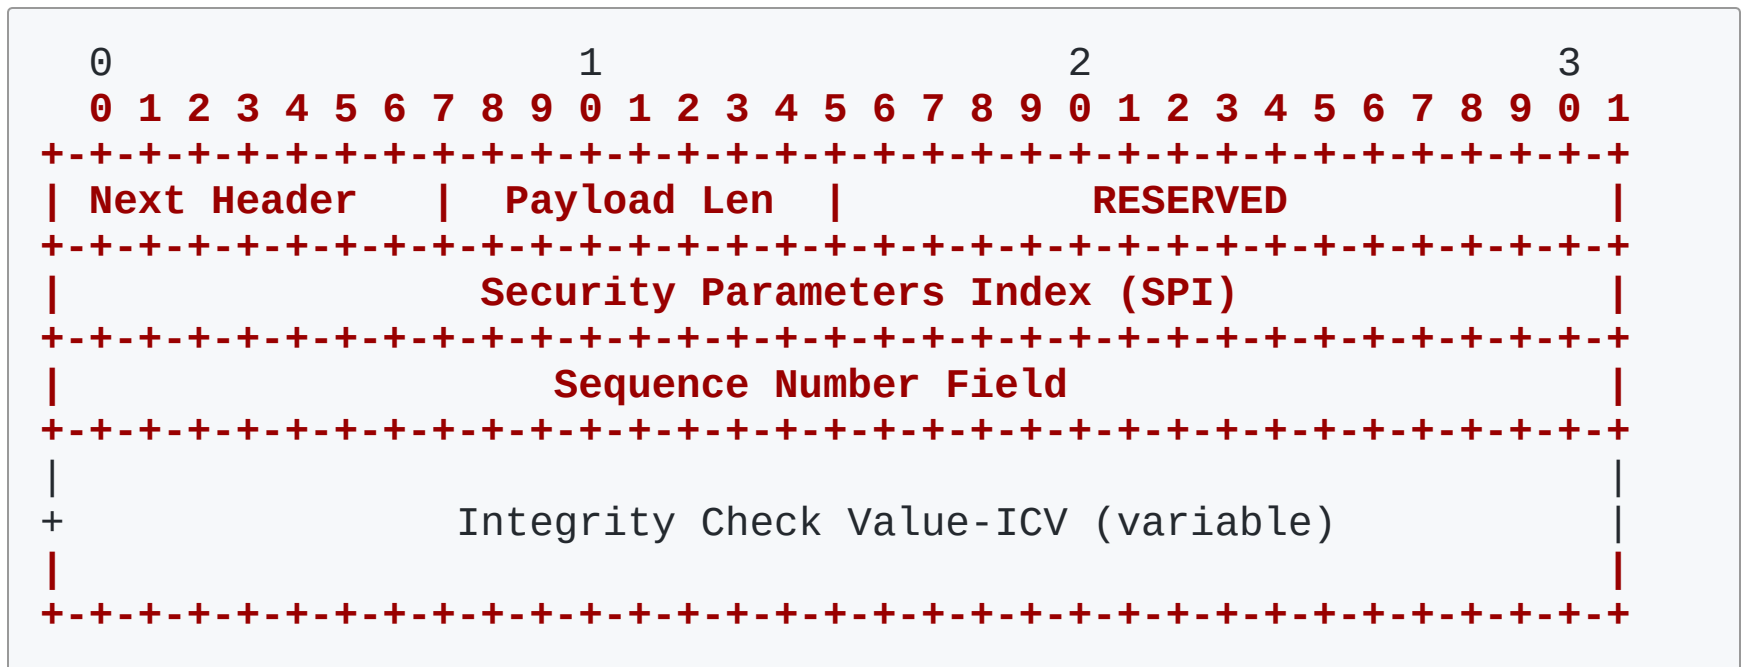
IPv6

```
-----  
| new* | new ext |          | orig* | orig ext |          |          | ESP  | ESP |  
| IP hdr| hdrs*  | ESP | IP hdr| hdrs *  | TCP | Data | Trailer | ICV |  
|          |          |          |          |          |          |          |          |  
-----
```

```
          |<----- encryption ----->|  
          |<----- integrity ----->|
```

* = **if** present, construction of outer IP hdr/extensions **and** modification of inner IP hdr/extensions is discussed **in** the Security Architecture document.

IPsec Authentication Header



IPsec AH Transport mode

BEFORE APPLYING AH

IPv6	-----			
		ext hdrs		
	orig IP hdr	if present	TCP	Data

AFTER APPLYING AH

IPv6	-----						
		hop-by-hop, dest*,		dest			
	orig IP hdr	routing, fragment.	AH	opt*	TCP	Data	

<--- mutable field processing -->				<-- immutable fields -->			
<---- authenticated except for mutable fields ----->							

* = **if** present, could be before AH, after AH, **or** both

IPsec AH Tunnel mode

BEFORE APPLYING AH

```
IPv6 |-----|
| orig IP hdr | if present | TCP | Data |
|-----|
```

AFTER APPLYING AH

```
IPv6 |-----|
| new IP hdr* | if present | AH | orig IP hdr* | if present | TCP | Data |
|-----|
|<--- mutable field -->|<----- immutable fields ----->|
|      processing      |
|<-- authenticated except for mutable fields in new IP hdr -->|
```

* = **if present**, construction of outer IP hdr/extensions **and** modification of inner IP hdr/extensions is discussed **in** the Security Architecture document.

IKEv2

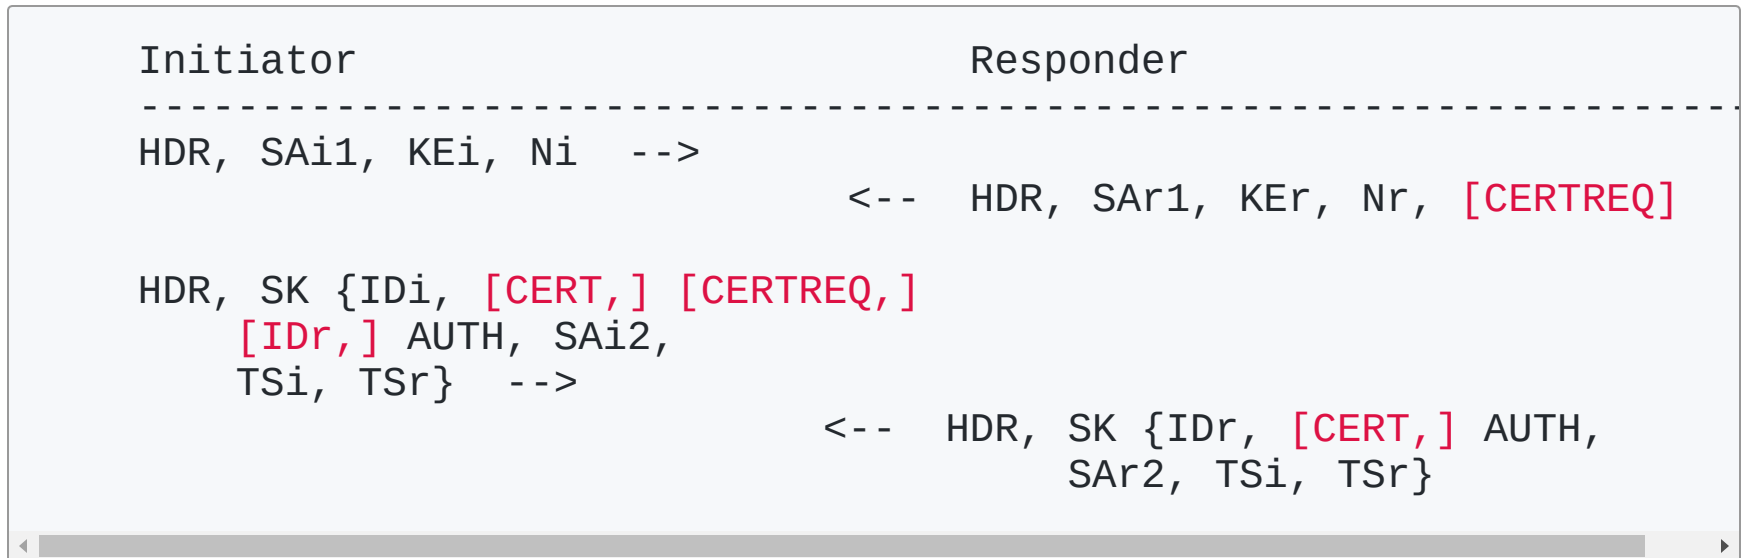
IKEv2 is the Internet Key Exchange Protocol:

- IKEv2 authenticates the peers
- Establishes a control channel between the peers
 - IKE_SA
- Negotiates the SA between the peers
- Manage the SA:
 - SA life cycle (rekeying)
 - Update SA in mobility / multihoming scenarios

IKEv2

The IKEv2 negotiation performs the following exchange type:

- IKE_INIT
- IKE_AUTH
- CREATE_CHILD_SA



INFORMATIONAL exchanges are performed latter

IKEv2 IKE_INIT

IKE_INIT results in the agreement of:

- SKEYSEED from which all IKE_SA keys are derives
- cryptographic algorithm use for the IKE channel (IKE_SA)

Initiator

Responder

HDR, SAI1, KEi, Ni -->

<-- HDR, SAR1, KEr, Nr, [CERTREQ]

- HDR: IKE Header - SPI, version, Exchange Type, Message ID
- SAI1, SAR1: proposal and selection of IKE_SA crypto algorithm
- KEi, KEr: public Diffie Hellman Value to generate SKEYSEED
- Ni, Nr: nonces to generate SKEYSEED

At this stage the peers have exchanged a shared secret, but...

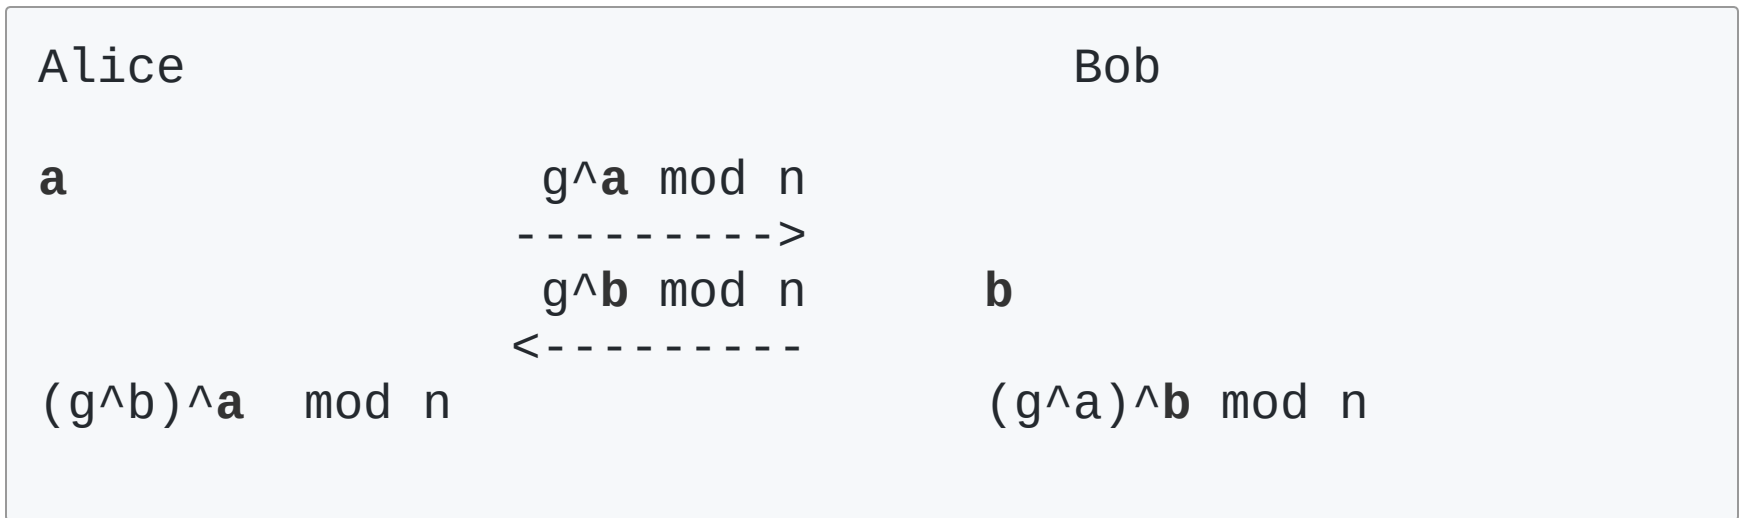
- Peers are still unauthenticated
- The IKEv2 channel is still unencrypted
- Negotiation of the SA has not yet even started.

Focus on Diffie Hellman

Diffie Hellman enables Alice and Bob to agree on a shared secret without revealing that secret.

Alice and Bob agree on a finite cyclic group G :

- of order n -- a very large number (2048 bit long)
- and a generator g in G .



Discrete logarithm says that it is hard to infer a from g^a

Focus on Key derivation

Protected communications usually involve multiple keys:

- Each key is usually dedicated to a specific task:
 - encryption, authentication, ...
 - initiator, responder, ...

To avoid one negotiation per key, we usually have:

1. Negotiation of a shared secret
2. Derivation of the keys from the share secret

Focus on Key derivation

The agreed shared secret is used to derive multiple secrets to:

- Protect the IKE_SA IKEv2 Channel:
 - Authentication SK_a
 - Encryption SK_e
- Derive the keying material associated to the SA
 - SK_d
- Authenticate the peers
 - SK_p

```
SKEYSEED = prf(Ni | Nr, gir)
{SK_d | SK_ai | SK_ar | SK_ei | SK_er | SK_pi | SK_pr}
    = prf+ (SKEYSEED, Ni | Nr | SPIi | SPIr)
```


IKEv2 IKE_AUTH

Initiator

Responder

HDR, SK {IDi, [CERT,] [CERTREQ,]
[IDr,] AUTH } -->

<-- HDR, SK {IDr, [CERT,] AUTH}

- SK: encryption key
- IDi, IDr: Identities
- CERT: public key associated to ID (optional)
- CERTREQ: Trusted Anchors
- AUTH: authentication payload

Focus on Digital Signature

The purpose of digital signature is to cryptographically proves:

- It is really *me* sending the message
- The message you receive is the one I have sent, e.g. has not been altered

A cryptographic signature is:

1. Generated by the owner of the private key
2. Validated with the associated public key, the signature and message

Alice

```
k_prv = private_key,  
k_pbl = public_key,  
m = message
```

```
s = sign(k_prv, m)
```

m, s, k_pbl

----->

verify(m, s, k_pbl)

Bob

Focus on AUTH

AUTH is not limited to a single message m , but the IKE_SA, that is:

- IKE_INIT (where the shared secret has been agreed)
- IKE_AUTH with IDs, CERT

AUTH can be seen as the signature where the signed data is:

- A known structure with sent and received information
- Built by each peer for signing / verifying

What has been sent corresponds to what has been received

Focus on AUTH

Building AUTH:

- build data
- sign data

Verifying AUTH:

- build peer's data
- verifying the signature

```
AUTHi = sign( IKE_INIT_request', Nr, prf(SKpi, IDi) )  
AUTHr = sign( IKE_INIT_response', Ni, prf(SKpr, IDr) )
```

Focus on IDs

AUTH enables to binds the IKE_SA to the owner of the private key

This owner is designated by its ID, but:

- How to ensure ID owns the private key ?
- How to make sure it is an authorized peer ?

CERT binds the ID and the public Key

- This may involve a Trusted Anchor (CA)

PAD indicates the peer is authorized

IKEv2 CREATE_CHILD_SA

CREATE_CHILD_SA negotiates the SA:

- Traffic Selectors (TS)
- Cryptographic Keys and algorithms
- Protocols (ESP, AH)
- IPsec Mode (default Tunnel)
- ...

IKEv2 CREATE_CHILD_SA

Initiator

Responder

HDR, SK {SAi2, TSi, TSr} -->

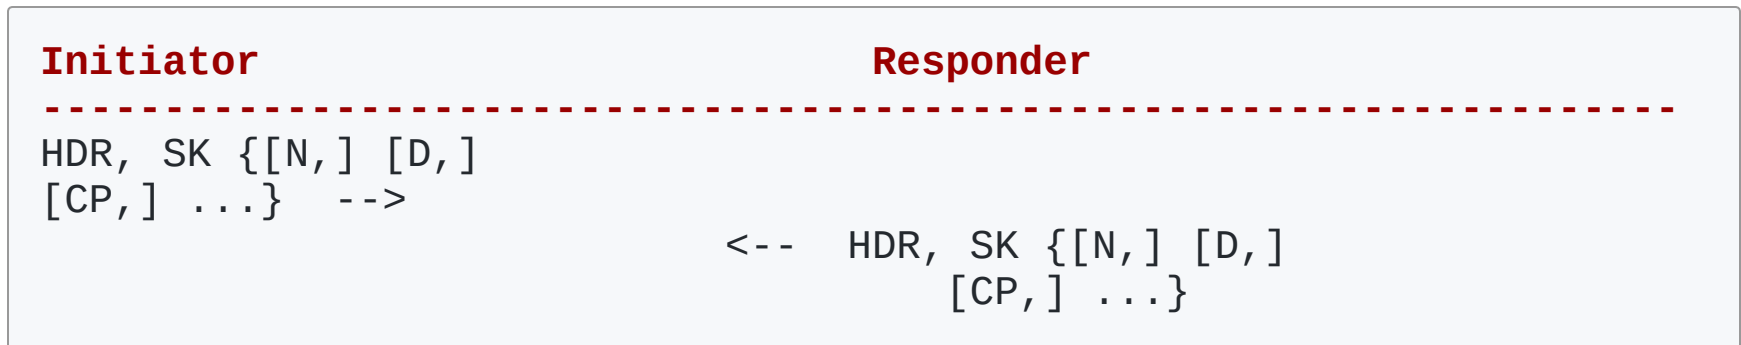
<-- HDR, SK {SAr2, TSi, TSr}

- SAI2, SAr2: proposal and selection of IKE_SA crypto algorithm
- TSi, TSr: Traffic Selectors

IKEv2 INFORMATIONAL exchanges

INFORMATIONAL exchange are based on query / response.

- used to Delete SA among other



- D Delete Payload

Other exchange mais involve other payloads:

- N Notify Payload
- CP : Configuration Payload (VPN)
- ...

VPN Example: Initial SPD - Security Gateway

SPD of the VPN Security Gateway

- BYPASSing IKEv2 channel

Traffic Selectors	PFP	SPD Cache	Processing
IP_dst: IP_{MN}^{outer}	0	IP_{MN}^{outer}	Policy: BYPASS
IP_src: IP_{SG}	0	IP_{SG}	
Port_dst: 500	0	500	
Port_src: 500	0	500	

Traffic Selectors	PFP	SPD Cache	Processing
IP_dst: IP_{SG}	0	IP_{SG}	Policy: BYPASS
IP_src: IP_{MN}^{outer}	0	IP_{MN}^{outer}	
Port_dst: 500	0	500	
Port_src: 500	0	500	

VPN Example: Initial SPD - Security Gateway

- Protecting the VPN Client traffic

Traffic Selectors	PFP	SPD Cache	Processing
Name: company.com	0	-	Policy: PROTECT
IP_dst: ANY	0	ANY	Mode: Tunnel
IP_src: ANY	1	IP_{MN}^{inner}	Tunnel Header: IP_{MN}^{outer} , IP_{SG}
Port_dst: ANY	0	ANY	Protocol: ESP
Port_src: ANY	0	ANY	Algorithm: AES-GCM

Traffic Selectors	PFP	SPD Cache	Processing
Name: company.com	0	-	Policy: PROTECT
IP_dst: ANY	1	IP_{MN}^{inner}	Tunnel Header: IP_{SG} , IP_{MN}^{outer}
IP_src: ANY	0	ANY	Mode: Tunnel
Port_dst: ANY	0	ANY	Protocol: ESP
Port_src: ANY	0	ANY	Algorithm: AES-GCM

VPN Example: Initial SPD, SAD - Security Gateway

- Preventing any other traffic

Traffic Selectors	PFP	SPD Cache	Processing
IP_dst: ANY	0	ANY	Policy: DISCARD
IP_src: ANY	0	ANY	
Port_dst: ANY	0	ANY	
Port_src: ANY	0	ANY	

- SPD Cache has not been populated
- SAD has not been populated

VPN Example: Initial SPD, SAD - VPN Client

SPD of the VPN Client

Traffic Selectors	PFP	SPD Cache	Processing
Name: IKEv2 UID	0	-	Policy: PROTECT
IP_dst: ANY	0	ANY	Mode: Tunnel
IP_src: ANY	1	IP_{MN}^{inner}	Tunnel Header: IP_{MN}^{outer} , IP_{SG}
Port_dst: ANY	0	ANY	Protocol: ESP
Port_src: ANY	0	ANY	Algorithm: AES-GCM

Traffic Selectors	PFP	SPD Cache	Processing
Name: IKEv2 UID	0	-	Policy: PROTECT
IP_dst: ANY	1	IP_{MN}^{inner}	Tunnel Header: IP_{SG} , IP_{MN}^{outer}
IP_src: ANY	0	ANY	Mode: Tunnel
Port_dst: ANY	0	ANY	Protocol: ESP
Port_src: ANY	0	ANY	Algorithm: AES-GCM

VPN Example: Initial SPD, SAD - VPN Client

- VPN Client BYPASS and DISCARD Policies are as those of Security Gateway
- SP has not been populated in the SPD Cache
- SA has not been populated in the SAD

VPN Example: IKEv2 Exchange

VPN Client initiates an Session with Security Gateway

- Exchange happen between $IP_{MN}^{outer}:500$ and $IP_{SG}:500$

```
Initiator                                Responder
-----                                -
HDR, SAI1, KEi, Ni  -->
                                <-- HDR, SAR1, KEr, Nr

HDR, SK { IDi, CERT, AUTH,
          CP(CFG_REQUEST),
          SAI2, TSi, TSr }  -->
                                <-- (IP_R1:500 -> IP_I1:500)
                                   HDR, SK { IDr, CERT, AUTH,
                                             CP(CFG_REPLY),
                                             SAR2, TSi, TSr }
```

- IPsec mode is tunnel
 - Tunnel IPs (IP_{MN}^{outer} , IP_{SG}) are read from IP header
- IP_{MN}^{inner} is provided by the Configuration Payload CFG_REPLY

VPN Example: SPD Cache

SPD of the VPN Security Gateway

- BYPASSing IKEv2 channel

Traffic Selectors	Processing
IP_dst: IP_{MN}^{outer}	Policy: BYPASS
IP_src: IP_{SG}	
Port_dst: 500	
Port_src: 500	

Traffic Selectors	Processing
IP_dst: IP_{SG}	Policy: BYPASS
IP_src: IP_{MN}^{outer}	
Port_dst: 500	
Port_src: 500	

VPN Example: SPD Cache

- Protecting the VPN Client traffic

Traffic Selectors	Processing
IP_dst: ANY	$SA_{Cl t \rightarrow SG}$
IP_src: IP_{MN}^{inner}	
Port_dst: ANY	
Port_src: ANY	

Traffic Selectors	Processing
IP_dst: IP_{MN}^{inner}	$SA_{SG \rightarrow Cl t}$
IP_src: ANY	
Port_dst: ANY	
Port_src: ANY	

VPN Example: SPD Cache

- Preventing any other traffic

Traffic Selectors	SPD Cache	Processing
IP_dst: ANY	ANY	Policy: DISCARD
IP_src: ANY	ANY	
Port_dst: ANY	ANY	
Port_src: ANY	ANY	

VPN Example: SAD

SAD of the VPN Client, Security Gateway

Action	Traffic Selectors	IPsec Signaling	Crypto material
$SPI_{Cl t \rightarrow SG}$	IP_dst: ANY	$IP_{src}^{Tunnel}: IP_{MN}^{outer}$	Crypto: K_e, K_a
	IP_src: IP_{MN}^{inner}	$IP_{dst}^{Tunnel}: IP_{SG}$	Counters: C_{ESP}, C_w
	Port_dst: ANY	Proto: ESP	
	Port_src: ANY	Algo: AES-GCM	
$SPI_{SG \rightarrow Cl t}$	IP_dst: IP_{MN}^{inner}	$IP_{src}^{Tunnel}: IP_{SG}$	Crypto: K_e, K_a
	IP_src: ANY	$IP_{dst}^{Tunnel}: IP_{MN}^{outer}$	Counters: C_{ESP}, C_w
	Port_dst: ANY	Proto: ESP	
	Port_src: ANY	Algo: AES-GCM	

VPN Example: Encrypted Traffic

```

0      1      2      3
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-----+-----+-----+-----+-----+-----+-----+-----+
I|version| traffic class |               flow label               |
P+-----+-----+-----+-----+-----+-----+-----+-----+
v|               payload length       | next header | hop limit |
6+-----+-----+-----+-----+-----+-----+-----+-----+
~                               inner source IP                               ~
+-----+-----+-----+-----+-----+-----+-----+-----+
~                               inner destination IP                               ~
+-----+-----+-----+-----+-----+-----+-----+-----+
E|               Security Parameters Index (SPI)               | ^
S+-----+-----+-----+-----+-----+-----+-----+-----+
P|               Sequence Number (SN)               |
+-----+-----+-----+-----+-----+-----+-----+-----+
|
|               IV               |
+-----+-----+-----+-----+-----+-----+-----+-----+
I|version| traffic class |               flow label               | ^
P+-----+-----+-----+-----+-----+-----+-----+-----+
v|               payload length       | next header | hop limit |
6+-----+-----+-----+-----+-----+-----+-----+-----+
~                               inner source IP                               ~
+-----+-----+-----+-----+-----+-----+-----+-----+
~                               inner destination IP                               ~
+-----+-----+-----+-----+-----+-----+-----+-----+
T|               source port         |               dest port         | r n
C+-----+-----+-----+-----+-----+-----+-----+-----+ y t
P|               Sequence Number (SN)               | p i
+-----+-----+-----+-----+-----+-----+-----+-----+ t c
|               ACK Sequence Number               | e a
+-----+-----+-----+-----+-----+-----+-----+-----+ d t
|Off. | Rserv |   Flags   |               Window Size               | e
+-----+-----+-----+-----+-----+-----+-----+-----+ d
|               Checksum               |               Urgent Pointer   |
+-----+-----+-----+-----+-----+-----+-----+-----+
|
~                               APPLICATION DATA                               ~
|
-|
E|
S+-----+-----+-----+-----+-----+-----+-----+-----+
P|               Padding (continue)       | Pad Length | Next Header | V V
+-----+-----+-----+-----+-----+-----+-----+-----+
|               Integrity Check Value-ICV   (variable)               |
+-----+-----+-----+-----+-----+-----+-----+-----+

```

MOBIKE

IP addresses of a packet are important Traffic Selectors likely used in:

- SPD, SPD Cache
- SAD

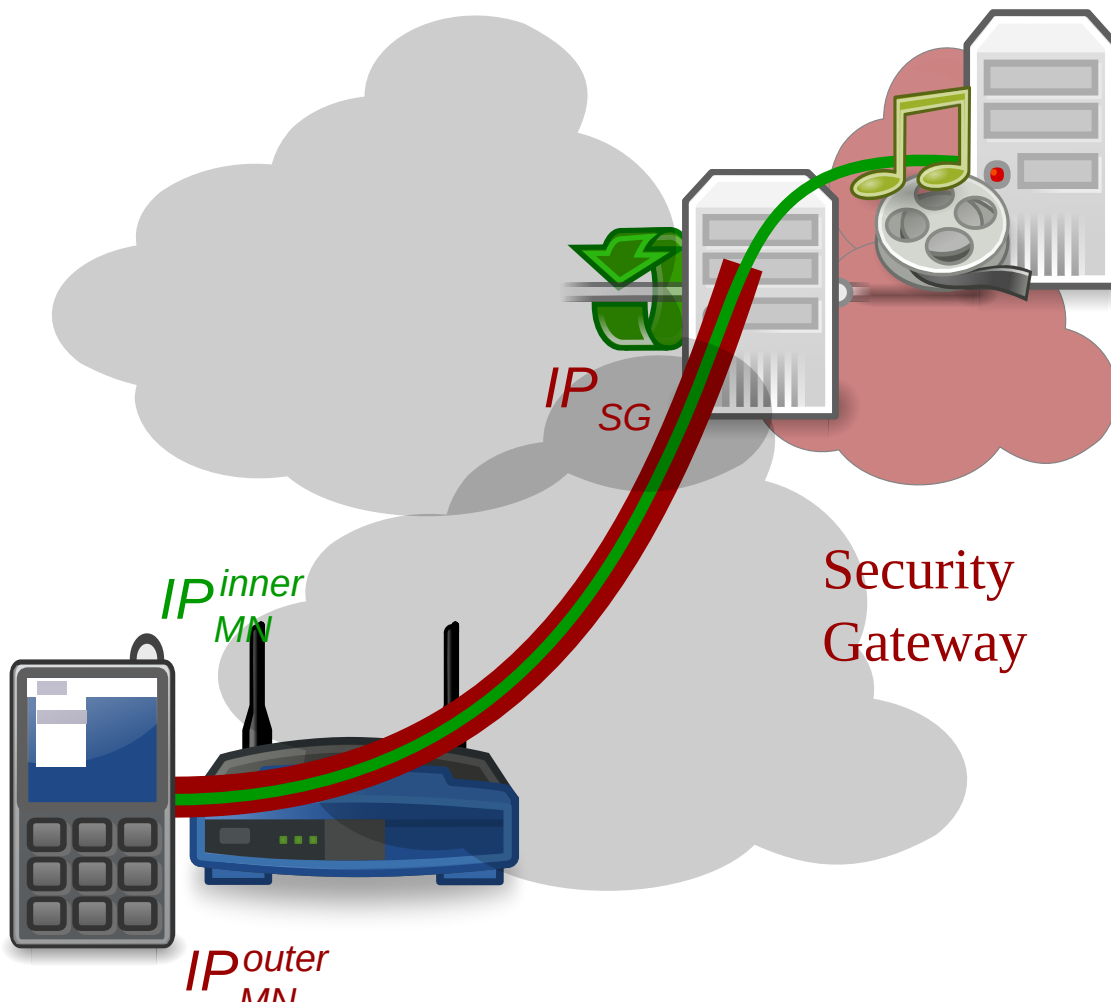
A Mobile Node (MN) IP changes need to be reflected in IPsec SAD/SPD:

- MOBIKE is the IKEv2 extension to enable updating IPsec configuration in Mobility and Multihoming environment

MOBIKE - Mobility

Initial Phase:

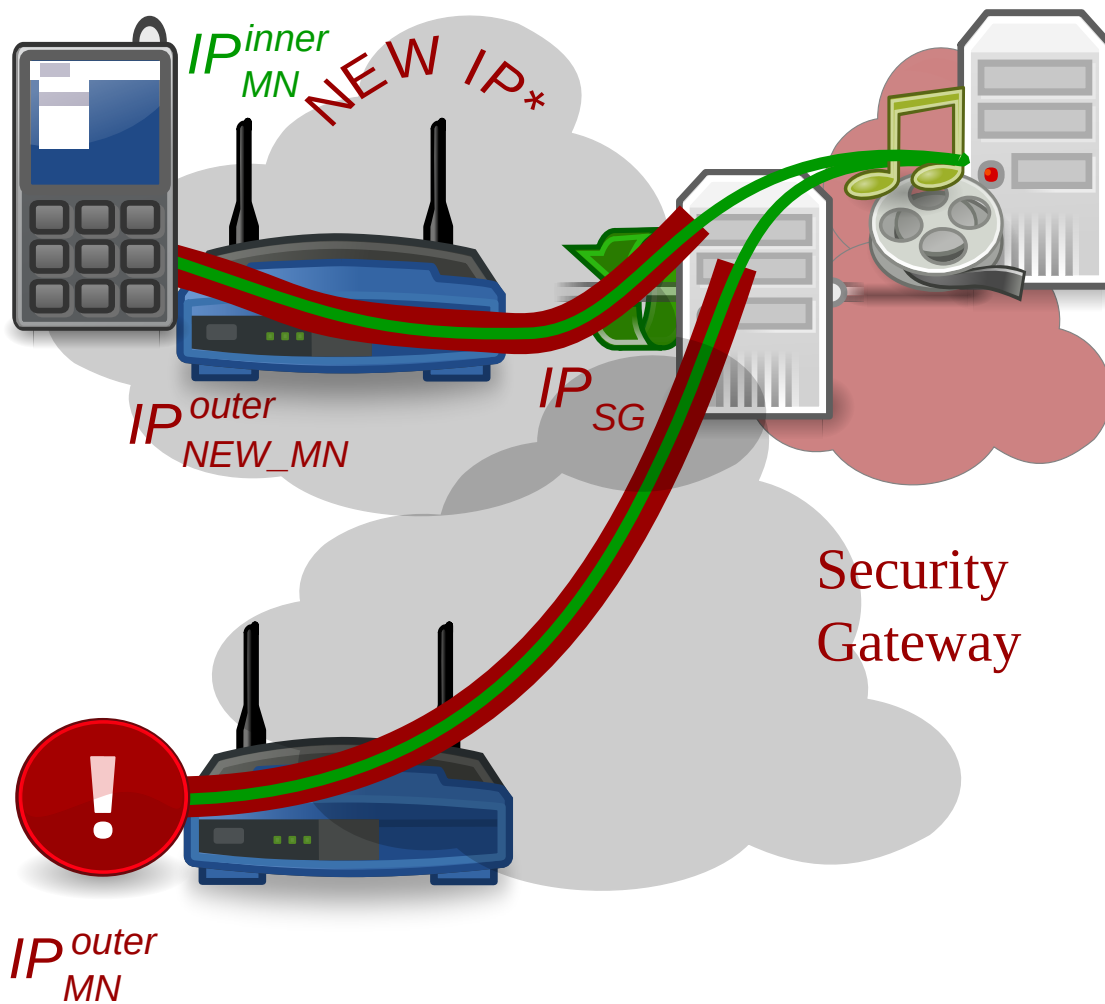
- MN connects a server with a private session: IP_{MN}^{inner} , IP_{MEDIA}^{inner}
- MN tunnels the private session to the Security Gateway: IP_{MN}^{outer} and IP_{SG}



MOBIKE - Mobility

Mobility Phase:

- MN changes ISP and the tunnel endpoint becomes IP_{MN}^{outer}



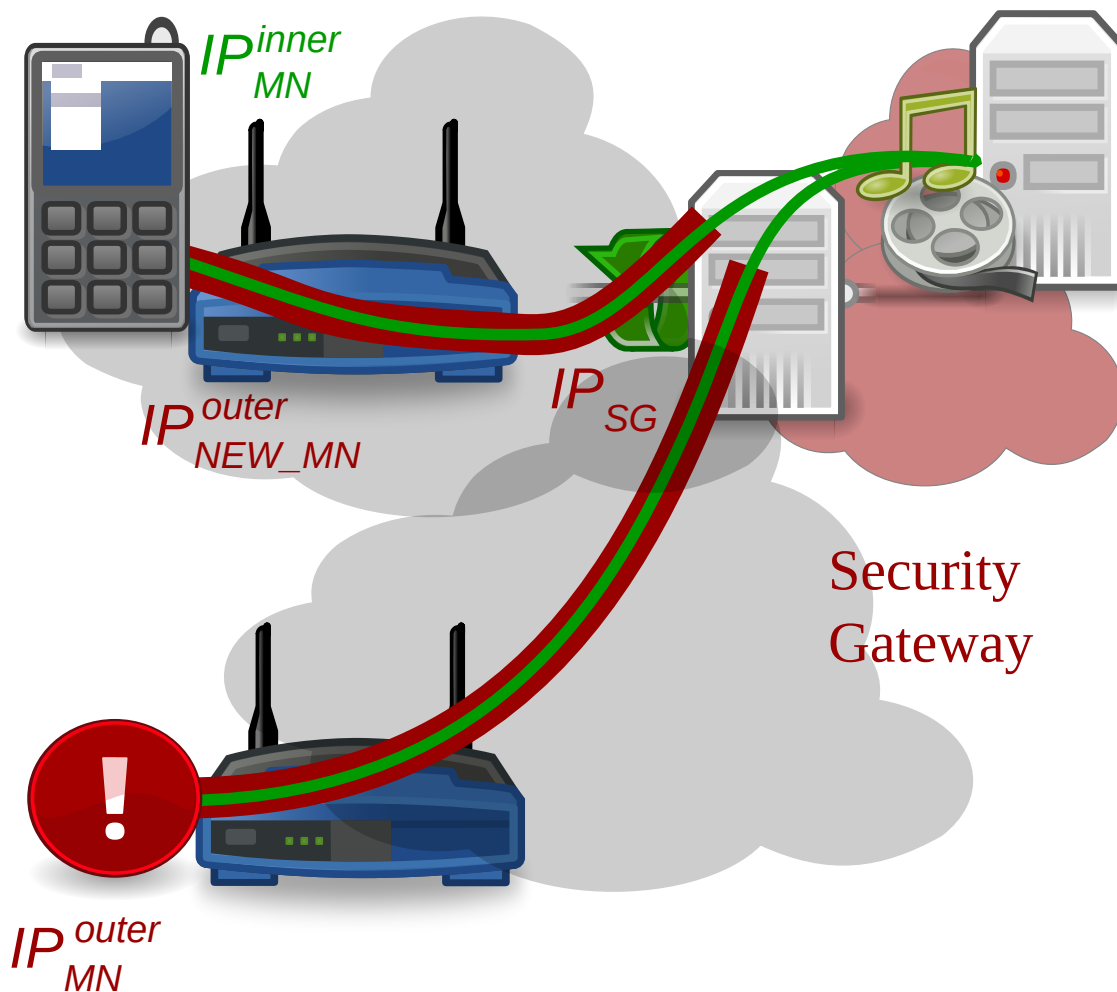
with advertisers using also reachable on $\mathbb{Z} = NEW_{MN} \cap \mathbb{Z} = MN$ items



MOBIKE - Multihoming

Mobility Phase:

- MN tunnel endpoint IP_{MN}^{outer} is not reachable
- Media server switches traffic to $IP_{NEW_MN}^{outer}$



MOBIKE - SPD - Cache

SPD Cache is not impacted by the change of IP address

- SPD concerns the inner traffic

Traffic Selectors	Processing
IP_dst: ANY	$SA_{Cl t \rightarrow SG}$
IP_src: IP_{MN}^{inner}	
Port_dst: ANY	
Port_src: ANY	

Traffic Selectors	Processing
IP_dst: IP_{MN}^{inner}	$SA_{SG \rightarrow Cl t}$
IP_src: ANY	
Port_dst: ANY	
Port_src: ANY	

MOBIKE - SAD Initial State

SAD (Outbound Traffic)

SAD of the VPN Client, Security Gateway

Action	Traffic Selectors	IPsec Signaling	Crypto material
$SPI_{Cl t \rightarrow SG}$	IP_dst: ANY	$IP_{src}^{Tunnel}: IP_{MN}^{outer}$	Crypto: K_e, K_a
	IP_src: IP_{MN}^{inner}	$IP_{dst}^{Tunnel}: IP_{SG}$	Counters: C_{ESP}, C_w
	Port_dst: ANY	Proto: ESP	
	Port_src: ANY	Algo: AES-GCM	
$SPI_{SG \rightarrow Cl t}$	IP_dst: IP_{MN}^{inner}	$IP_{src}^{Tunnel}: IP_{SG}$	Crypto: K_e, K_a
	IP_src: ANY	$IP_{dst}^{Tunnel}: IP_{MN}^{outer}$	Counters: C_{ESP}, C_w
	Port_dst: ANY	Proto: ESP	
	Port_src: ANY	Algo: AES-GCM	

SAD (Outbound Traffic)

SAD of the VPN Client, Security Gateway

Action	Traffic Selectors	IPsec Signaling	Crypto material
$SPI_{Cl t \rightarrow SG}$	IP_dst: ANY	$IP_{src}^{Tunnel}:$ $IP_{NEW_{MN}}^{outer}$	Crypto: K_e, K_a
	IP_src: IP_{MN}^{inner}	$IP_{dst}^{Tunnel}: IP_{SG}$	Counters: C_{ESP}, C_w
	Port_dst: ANY	Proto: ESP	
	Port_src: ANY	Algo: AES-GCM	
$SPI_{SG \rightarrow Cl t}$	IP_dst: IP_{MN}^{inner}	$IP_{src}^{Tunnel}: IP_{SG}$	Crypto: K_e, K_a
	IP_src: ANY	$IP_{dst}^{Tunnel}:$ $IP_{NEW_{MN}}^{outer}$	Counters: C_{ESP}, C_w
	Port_dst: ANY	Proto: ESP	

MOBIKE - IKEv2 Exchanges

In order to provide IPsec in mobility / multihoming environment

- SA needs to be updated in sync
- Mobility:
 - the MN needs to be able to update its SA
 - the MN needs to inform the SG of the new IP address
 - the SG needs to be able to update the SA
- Multihoming:
 - the MN needs to provide alternate IP addresses
 - the SG needs to update its SA
 - the SG needs to check reachability
 - the MN needs to update its SA

MOBIKE - Mobility - IKEv2 Initial Exchanges

Initiator	Responder
-----	-----
1) (IP_I1:500 -> IP_R1:500) HDR, SAI1, KEi, Ni, N(NAT_DETECTION_SOURCE_IP), N(NAT_DETECTION_DESTINATION_IP) -->	<-- (IP_R1:500 -> IP_I1:500) HDR, SAR1, KER, Nr, N(NAT_DETECTION_SOURCE_IP), N(NAT_DETECTION_DESTINATION_IP)
2) (IP_I1:4500 -> IP_R1:4500) HDR, SK { IDi, CERT, AUTH, CP(CFG_REQUEST), SAI2, TSi, TSr, N(MOBIKE_SUPPORTED) } -->	<-- (IP_R1:4500 -> IP_I1:4500) HDR, SK { IDr, CERT, AUTH, CP(CFG_REPLY), SAR2, TSi, TSr, N(MOBIKE_SUPPORTED) }

MOBIKE - Mobility - Routability Check

Responder verifies that the initiator has given it a correct IP address.

```
4)                                <-- (IP_R1:4500 -> IP_I2:4500)
                                HDR, SK { N(COOKIE2) }

(IP_I2:4500 -> IP_R1:4500)
HDR, SK { N(COOKIE2) }  -->
```


MOBIKE - Multihoming - IKEv2 Initial Exchanges

Initiator	Responder
-----	-----
1) (IP_I1:500 -> IP_R1:500) HDR, SAI1, KEi, Ni, N(NAT_DETECTION_SOURCE_IP), N(NAT_DETECTION_DESTINATION_IP) -->	<-- (IP_R1:500 -> IP_I1:500) HDR, SAR1, KEr, Nr, N(NAT_DETECTION_SOURCE_IP), N(NAT_DETECTION_DESTINATION_IP)
2) (IP_I1:4500 -> IP_R1:4500) HDR, SK { IDi, CERT, AUTH, CP(CFG_REQUEST), SAI2, TSi, TSr, N(MOBIKE_SUPPORTED) } -->	<-- (IP_R1:4500 -> IP_I1:4500) HDR, SK { IDr, CERT, AUTH, CP(CFG_REPLY), SAR2, TSi, TSr, N(MOBIKE_SUPPORTED), N(ADDITIONAL_IP4_ADDRESS) }

MOBIKE - Detection

The initiator suspects a problem in the currently used address pair and probes its liveness.

```
3) (IP_I1:4500 -> IP_R1:4500)
   HDR, SK { N(NAT_DETECTION_SOURCE_IP),
             N(NAT_DETECTION_DESTINATION_IP) } -->

(IP_I1:4500 -> IP_R1:4500)
HDR, SK { N(NAT_DETECTION_SOURCE_IP),
          N(NAT_DETECTION_DESTINATION_IP) } -->
```

MOBIKE - Detection

Eventually, the initiator gives up on the current address pair and tests the other available address pair.

```
4) (IP_I1:4500 -> IP_R2:4500)
   HDR, SK { N(NAT_DETECTION_SOURCE_IP),
             N(NAT_DETECTION_DESTINATION_IP) }

      <- - (IP_R2:4500 -> IP_I1:4500)
           HDR, SK { N(NAT_DETECTION_SOURCE_IP),
                     N(NAT_DETECTION_DESTINATION_IP)
```

MOBIKE - Update

This worked, and the initiator requests the peer to switch to new addresses.

```
5) (IP_I1:4500 -> IP_R2:4500)
   HDR, SK { N(UPDATE_SA_ADDRESSES),
             N(NAT_DETECTION_SOURCE_IP),
             N(NAT_DETECTION_DESTINATION_IP),
             N(COOKIE2) } -->

               <-- (IP_R2:4500 -> IP_I1:4500)
                   HDR, SK { N(NAT_DETECTION_SOURCE_IP),
                             N(NAT_DETECTION_DESTINATION_IP),
                             N(COOKIE2) }
```

Reference

- Kent, S. and K. Seo, "Security Architecture for the Internet Protocol", [RFC 4301](#), December 2005
- Kent, S. "IP Authentication Header", [RFC 4302](#)
- Kent, S., "IP Encapsulating Security Payload (ESP)", [RFC 4303](#), December 2005
- Kaufman, C., Hoffman, P., Nir, Y., Eronen, P., and T. Kivinen, "Internet Key Exchange Protocol Version 2 (IKEv2)", [RFC 7296](#), October 2014
- P. Eronen, "IKEv2 Mobility and Multihoming Protocol (MOBIKE)" [RFC 4555](#), June 2006

Acknowledgment

- Open Security Architecture